

Web search engines work by storing information about a large number of Web pages. These pages are retrieved by a crawler, also known as a spider. The content of every page is then analysed to determine its indexing. Some search engines, such as Google, store all or part of the source page (referred to as a cache) as well as information about the Web pages. When a user submits a query by giving keywords, the engine looks up the index and provides a listing of best-matching Web pages according to its criteria.

The usefulness of a search engine depends on the relevance of the results it gives. While there may be millions of Web pages that include a particular word or phrase, some pages may be more relevant, popular, or authoritative than others. Most search engines employ methods to rank the results to provide the 'best' results first. How a search engine decides which pages are the best matches, and what order the results should be shown in, varies widely from one engine to another. The methods also change over time as Internet usage changes and new techniques evolve.