

tion or device. You cannot make files of one partition appear as part of another. This layout, called a *unified filesystem*, allows files from one drive to appear as part of another. Hence you need to **mount** the partition you wish to access, to a mount point. A mount point is the directory to which the partition is attached. After mounting, the contents of the partition appear in the mount point.

What is the use of this method of working?

Let's say that in your Windows "D:" partition you have some songs, and in partition "E:", you've copied songs from a friend. You want to know what songs are common. You can simply mount both these partitions to a single mount point in Linux and view them under it! Arrange files according to name, and your task is done. This is a trivial example, but the method can be of great use at the network administration level. Scripts that operate on files in a directory can be used on files of different partitions by simply mounting them to a common mount point.

The following codes represent hard disks connected in different configurations:

Primary Master—First IDE hard disk—hda

Primary Slave—Second IDE hard disk—hdb

Secondary Master—Third IDE hard disk—hdc

Secondary Slave—Fourth IDE hard disk—hdd

Further, the first partition on a disk is represented by 0, the second partition by 1 and so on.

Hence, hdb2 refers to the third partition on the second IDE device (primary slave).

If you don't know which partition you wish to mount, you can get the partition information of all connected hard disks via the "**fdisk -l**" command and then proceed to mount.

The mount command is used to mount the partitions into Linux. The general format of the command is: