

Duane Wessels of the National Laboratory for Applied Network Research. Being freeware, it wasn't commercially interesting enough to keep the money coming in, and a few years later it ran out of funds. Nevertheless, the project continued through donations and occasional commercial investments (a list of sponsors is included in the download package!)

Squid was originally designed to run on Linux / UNIX systems, and the first beta release (v1.2) came out in October 1997. Being a constantly-evolving project, bugs were fixed and new features added every now and then by a group of people devoted to spending their time and effort developing and enhancing it, until the first "stable" version (v2.3) came out in January 2000.

Five years and two releases later, the first Windows version (v2.5) came out in May 2005. The current release, v2.6, is by far the latest and most stable available.

12.5.2 Why Squid It?

Or a more typical question; why do I need to cache anything at all? Why can't my browser just get me my page fresh off the Web instead of dumping it to disk and running several risks—from getting a private page cached to obtaining stale versions of pages when requested?

Well, this wouldn't happen for two reasons. First, Squid is based on a strict protocol that defines exactly what can be cached and what shouldn't, how long would a page remain fresh, and when to obtain the page from the Web server (and update the cache with this new information). So when it caches information from frequently-visited Web sites and builds up a repository of stored data, it knows exactly what to cache and how long to keep it.

Secondly, Squid has an advanced content routing configuration that allows you to build content clusters to route requests and balance loads. In plain English, that means Squid would service