

requests by just copying the content being used instead of copying the whole thing all over again. So even if the requested information is available in the cache, Squid verifies whether that data is still fresh before sending it over to the client. If its not, then the part that isn't fresh is requested from the Web server, merged with the existing data, and sent to the client.

So adding up the numbers, in a large network it makes a lot of sense to store and reuse cached data instead of giving fresh requests to the same (unchanged) page that has already been retrieved hundreds of times over the past few days. And oh, Squid is free, so anyone can download and use it.

12.6 Getting Started

12.6.1 The Basics

Like we mentioned, Squid was originally designed to be Linux / UNIX native, and so was not intended to be run on Windows-based servers up until release v2.5. But that's all in the past; Squid v2.6 currently supports the following servers:

- Windows NT4 (Service Pack 4 or later)
- Windows2000 (Server or Professional)
- Windows Server 2003
- Windows XP Professional

The Windows version can be downloaded free from the Squid homepage (see links at the end of this chapter), and you'll find the Linux version on this month's DVD.

Okay, so you're all excited about setting up Squid and got it downloaded. Before you open the ZIP file and start to wonder where the familiar icons and programs are, remember that since this was never *really* designed with Windows in mind, it doesn't feature the Windows-style console interface we're all used to.