

Each request is recorded in a standard log file format. Being standardised, log files are highly accurate criteria by which the popularity of a Web site could be established. Some of these log files are:

Access_log: The principal log file that records filenames, IP addresses, date, time, and other relevant data.

Referrer_log: Stores the URLs of the sites from which the user was referred to this site.

Error_log: Requests for files and error messages are entered into this file.

Perhaps a really simple way to find how many requests a server received is to count the number of lines in the log, as each log entry represents a request. Besides, each log line shows a request for a specific URL that comes from a definite IP address. This could make it even easier to trace how many requests a particular page on the Web site got and how many came from a single source.

2.2.4 Felling Logs

The way to analyse Web site traffic is to be able to understand what exactly a typical log entry in a Web server actually means. A typical log entry would look something like this:

```
123.123.45.13 - [15/Jan/2008:11:30:22 -0400] "GET/stuff.html  
HTTP/1.0" 200 16716
```

The data in the Access_log file above is made up of several different fields. A file in the Common Log Format typically has eight fields of information for each HTML or graphics file served. Going from left to right, the fields that are listed in the log are:

1. Address field (123.123.45.13): This is obviously an IP address. The address field records the domain name or IP address from which the request has been made. So if you are asking for the Yahoo! home page, the Web server would make a note of *your* IP address (not the IP address of the Yahoo! site) as the first item of data in the access log. Often, system administrators choose to store only the IP address and not the domain names, because that would require