

- Configure the number of instances you would need as well as their capacities. Please note that the underlying EC2 instances will be charged separately.
- Once selected, you will be asked to configure the cluster type. To do a map reduce data streaming job, select the streaming option and configure the location of the mapper and reducer. Set the location of the output jobs and logs, and you are good to go.
- Click the next button to launch the cluster, you will see metrics of the cluster including the running time, current status, as well as ARN. You can monitor the cluster state and once the mapping and reducing jobs are complete, the cluster should automatically terminate.



Point out the location of the mapper and the reducer.

Thats it! You have now successfully launched an EMR cluster that has run a MapReduce job over a distributed system and stored the results in an S3 bucket.

Redshift

Of course, all this processing and computation is meaningless unless you have a means of storing the big data in a way that allows you to access it as and when needed. Traditionally this was done either in file systems (ugh), or databases. Both of them however, need to be extensively tweaked to maintain speed and throughput in the face of the very high volume that big data can potentially generate. AWS can help you here too. The Redshift data warehouse service is a very fast and fully managed solution that scales to petabytes and makes it simple and cost-effective to efficiently analyse all your data without changing the choice of your Business Intelligence (BI) tools.

You can very easily setup your own redshift warehouse and start logging all your data into it.

Setting up AWS Redshift:

- Navigate to “Redshift” from the “Services” tab. On opening the Redshift console, you will be presented with a home screen that lists your