



Launch a new Redshift cluster to save all your processed data.

current warehouse clusters. If there aren't any, click on the “Launch Cluster” button.

- Fill in the details on the next screen such as the cluster identifier that uniquely identifies the redshift cluster, the ‘database name’ that can be used to connect to the redshift database, as well as the port that can be used for the connection (defaults to 5439). Set a master username and password.
- Next, set the node types (these vary in compute capacity and number of EC2 per node) and the number of nodes required for the cluster.
- Finally, on the next screen fill out the security details such as choice of VPC, security groups and availability zone etc. Click continue to deploy your cluster.

[Please Note: Deploying a cluster might incur high charges, unless you are eligible for the ‘free tier’. Please double check configuration before deploying.]

The cluster will take a few minutes to spin up. Once running, you can connect to it using a JDBC driver and the url `jdbc:redshift://<cluster-arn>:<port>/<database-name>`

Now you can start adding data to the cluster as well as connect all your BI tools to the redshift cluster for running analytics. **d**