

re-written as a set of data points for a program to use. An example of such a problem would be our ability to recognise faces. Even though infants are able to do it (to some degree at least) we find it hard to quantify how exactly we recognise a face we have seen before. What factors are at play when we are trying to do it and which face similarity qualifies as recognition. These problems when translated for computers require them to 'look' at images in a certain context and come up with answers. They differ from data related problems to some degree in the approach towards their solution. Amazon offers a set of services under its 'Rekognition' service to tackle machine learning problems that involve images. We will now take a look at how to go about building applications on it.

Amazon Rekognition allows you to just provide an image (or video) as input to the Rekognition API, and the service can then identify certain objects, people, text, scenes, and activities within it, as well as detect any inappropriate content. It also provides highly accurate facial analysis and facial recognition and can detect, and compare faces for a wide variety of use cases such as user verification, security, etc. Rekognition uses a highly scalable, deep learning technology developed by Amazon's computer vision scientists that is capable of analysing a huge number of images and videos daily, and just like Amazon ML requires almost no machine learning expertise to use. It is designed as a simple and easy to use API that can quickly analyse any image or video file stored in Amazon S3. Amazon Rekognition is always learning from continuous feedback data that is automatically leveraged for all future runs.

## Getting Started with Amazon Rekognition

- To start using the Amazon Rekognition API, you would need to have Amazon CLI installed on your machine. In case you don't have it, please install it before continuing further. The steps to install Amazon CLI have already been discussed in the introductory chapter.
- Once the AWS CLI is set up, try running the below command with a test image to ensure that Rekognition is working correctly.
- ```
>aws rekognition detect-labels --image "S3Object={Bucket=<test-bucket>,Name=example.jpg}" --region us-west-2
```

This would run Rekognition against an example image located at `s3://test-bucket/example.jpg` and return the labels of objects found in it (person, building etc).