


Amazon SageMaker allows directly deployment of machine learning models from training into a production-ready hosted environment. It provides an integrated Jupyter authoring notebook instance for easy access to your data sources for exploration and analysis, so that you do not have to have a separate server for it. It also provides common machine learning algorithms with native support for custom algorithms and frameworks. Amazon SageMaker also offers flexible distributed training options that allow you to autoscale your training jobs according to a specific workflow.

Getting Started on Amazon SageMaker:

- Navigate to 'Amazon SageMaker' under Services tab and click on 'Create new notebook instance' on the SageMaker console. This allows you to create a new notebook instance that can be used to design your machine learning model.
- Fill out the notebook instance name and assign a role to it and click continue to successfully create the notebook. Optionally you can also specify an encryption mode for it.



Create a new notebook instance on the SageMaker console.

- Once the notebook is created, you can choose to open it from the 'Notebook instances' tab list by clicking on the entry and then clicking the 'Open' button. This opens up the jupyter notebook where you can design the actual machine learning application.
- When the notebook is opened, click on the 'New' button and start writing out your algorithm in a platform of choice (spark, mxnet etc), you can