is anything but simple, but it seems like by the end of this decade, the deep learning industry will simplify its offerings so that they are comprehensible and useful to the average-Joe developer.

The deep learning industry will most probably converge on a set of standard tools, by the end of this decade, and will integrate them into all deep learning processes. The most popular tools as of now are TensorFlow, BigDL, OpenDeep, Caffe, Theano, Torch and MXNet. Another very real possibility is that deep learning with gain native support within the Spark community since they are already planning to beef up the platform's native deep learning capabilities.

Deep learning will possibly find its stable niche within the open analytics system. It is becoming abundantly clear that you cannot properly train, manage and deploy deep learning systems without the full suite of big data analytics capabilities which are provided by platforms like Spark, Hadoop, Kafka and more, therefore this is a very real possibility.

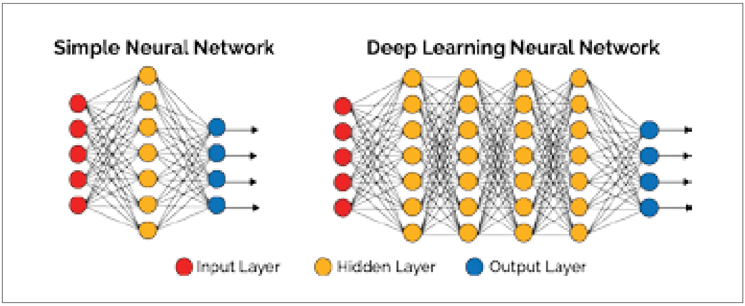

Deep learning may replace traditional machine learning in various areas.

We also predict that deep learning tools will be firmly implanted into every design surface, where the next decade will see deep learning development tools, libraries and languages becoming standard components in every software development kit. As the deep learning market takes strides towards mass adoption, chances are that it will follow the footsteps of sectors such as data visualisation, business intelligence and predictive analytics markets. All of these have moved their own solutions towards a self-service cloud-based delivery model that optimise user experience with fast delivery and also solutions that do not bother them with technical complexities. All-in-all, deep learning will be implemented in almost every imaginable industry since automation is the name of the game in the 21st century.

https://software.intel.com/en-us/