

Caffe and there's also a version that is optimised for Intel architecture that is currently integrated in Intel® Math Kernel Library (Intel® MKL) 2017. It's also optimised for Intel® Advanced Vector Extensions 2 (Intel® AVX2) and also Intel® Advanced Vector Extensions 512 (Intel® AVX-512). When the two versions were pitted against each other, the vanilla version of Caffe executed a project in 37.026 seconds, whereas, the optimised version completed the same project on the same hardware while only consuming 31.149 seconds. The 6 seconds of time saved might appear inconsequential but when it scales in context of an actual Deep Learning application, the difference is huge.

Intel has constantly highlighted the Intel® Xeon™ Scalable processor family as the most advanced yet boasting of several computational edges over its competitors. With Intel working on both the hardware and software offerings catered towards Deep Learning applications, it becomes a natural choice for many reasons not limited to compatibility, stability and optimisation. 