on a batch of images of 128 or 256 images at once in just a few milliseconds. However, the power consumption can double with more data.

3. Always check for three main metrics when looking for GPUs – memory bandwidth (how much data the GPU can handle), processing power (how quickly the GPU can compress data for each clock speed) and video RAM size (the amount of data on the video card at any point of time).
4. Using multiple video cards is a great boost for high level operations but should be restricted to 16 PCIe lanes available for data transfer. Processors for single desktops usually use 16 lanes. Anything that takes 32 lanes is not fit for CPUs. For 3 or 4 GPUs, use 8x the lanes per card.

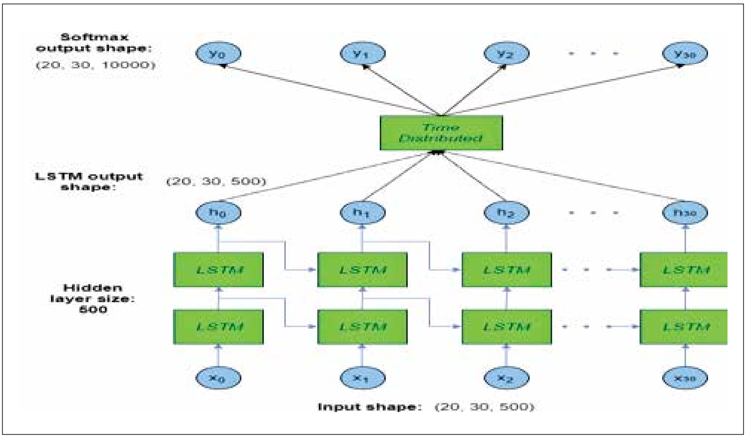

Keras Functional API LSTM Model

5. Just to avoid any start up issues use the CUDA drivers, which are discussed in the next chapter. Built by NVIDIA, the CUDA toolkit is well integrated with multiple platforms including Pytorch, sklearn, Tensorflow and Keras.
6. Intel recommends changing the parallelism threads and the OpenMP settings by changing the configuration files:-

```
•     config = tf.ConfigProto(intra _ op _ parallelism _ threads=NUM _ PARALLEL _ EXEC _ UNITS, inter _ op _
parallelism _ threads=2, allow _ soft _ placement=True,
device _ count = {'CPU': NUM _ PARALLEL _ EXEC _ UNITS })
session = tf.Session(config=config)
•     K.set _ session(session) os.environ["OMP _
```