




RELU Model for Keras

```
);cudaEventDestroy( start );cudaEventDestroy( stop );
```

The above code should be implemented whenever new events have to be removed and created in the CUDA mainframe.

```
• size=N*sizeof(float)/nStreams;for (i=0; i<nStreams; i++) {
•     offset = i*N/nStreams;     cudaMemcpyAsync(a_
d+offset, a_h+offset, size, dir, stream[i]);     kernel<<<N/
(nThreads*nStreams), nThreads, 0,
•     stream[i]>>>(a_d+offset);}
```

A look at other hardware

Experts suggest that the best libraries for performing deep learning on keras are available on NVIDIA which have taken over AMD's OpenCL and other alternatives. The difference in the years of release also means that you're likely to find better solutions for hardware issues for NVIDIA.

Next up in the list comes AMD. The company developed a format to convert the code into the respective GPU language before being compiled to the GPU assembly. Early users will, however, find the challenges while importing TensorFlow and PyTorch code bases from the mainframe. Deep learning support for Keras and TensorFlow exists for AMD's hardware but it is dogged with lots of coding and confusing syntaxes.

Performance wise, AMD GPUs now have 16-bit compute capability but still fail to match to the Tensor Cores of NVIDIA GPUs which work fantastically for transformers and convolutional networks.

And then comes Intel. The Xeon Phi has often been hailed to be the company's best hardware for deep learning projects but might not match up to the likes of NVIDIA or AMD cards. Newer patches and code com-