

that a large number of epochs will give higher accuracy, as seen in our second model. We discussed a bottom-up approach of using Keras for image recognition tasks - starting from pre-processing the data set, splitting it for training and validation, designing our model - even using the approach of transfer learning, compiling it and finally plotting various charts of loss and accuracy measures. We learnt about various models, layers and discussed the architecture of some of the recent well-known models such as the Inception, Xception, ResNet and VGG. Finally, we learnt how to gain insights from the plots regarding overfitting and the number of epochs required in increasing accuracy. **d**