

For classifier-free guidance, the model is trained with conditioning, but the caption is randomly dropped during training. This results in a model that can work both conditioned and unconditioned. During inference, the image is passed through the model with and without the conditioning caption, and the difference between the two is used to extrapolate the final prediction towards the direction of the conditioned prediction.

It's worth noting that a diffusion model can use both conditioning and guidance simultaneously. 