




Mahesh Benkar

The Unc

If a spammer sent you an e-mail with a subject line that said, "Your mail storage limit has changed," you'd probably open it, wouldn't you? You wouldn't know it was spam until you opened it... but CRM114—the Controllable Regex Mutilator—would. And it's 1,000 per cent more likely to classify the e-mail as spam than humans are. CRM114 is a spam killer that employs Bayesian Learning, a technique that's getting increasingly popular in spambusting.

Largely ignored so far, the possibilities afforded by Bayesian Belief Networks (BBNs) and Bayesian Learning (BL)—learning in BBNs—have only recently begun to be explored.

Spam detecting and search enhancing is just the tip of the iceberg when it comes to the possibilities of this technique, which, though it made its debut in 1984, had its basic theoretical foundation spelt out 250 years ago.

BL is a process whereby a machine is trained on past data so it can come to decisions about the future. The basic idea is that the likelihood of something happening can be estimated by how often it happened in the past. These decisions—that something will happen, or that something is true or false—are probabilistic and are