



Intel Lunar Lake

Bringing the fight to the frontlines

Soham Raninga | soham@digit.in

In a bold bid to reclaim its dominance, Intel has launched its most significant initiative to date: Lunar Lake. As Qualcomm and AMD gain ground, Intel is pulling out all the stops, making ambitious claims and putting its reputation on the line. At a press briefing in Berlin, Josh Newman, VP of Client Computing Group, set the stage for a paradigm shift in the AI PC landscape, stressing that a truly great AI PC demands a great overall PC. The Intel Lunar Lake series promises sweeping gains across the performance spectrum, boasting revamped CPU, GPU, and NPU engines optimized for efficiency and power. Key highlights include a significant leap in efficiency, directly challenging Qualcomm's battery life claims, and a dramatically improved integrated GPU performance courtesy of Intel XE2, aimed at amplifying AI capabilities and gaming performance. Let's get into the details of each performance engine and take a

closer look at how Intel plans to regain its leadership position on the performance and efficiency front.

Most power-efficient Intel CPU ever

Intel's Lunar Lake platform marks a significant milestone with the most power-efficient Intel CPU ever. The company has made bold and ambitious decisions, starting with 8-core configurations across the lineup, notably without hyperthreading. This deliberate choice aims to enhance performance per watt through a two-cluster approach, featuring four power-efficient E-cores for workloads that don't require high performance. Building on Meteor Lake's 2x Crestmont E-cores with a 2MB shared L2 cache, Lunar Lake boasts 4x Skymont E-cores with a 4MB shared L2 cache, introducing several enhancements. These include improved instruction prediction, increased parallelism with 26 dispatch ports, and deeper queueing for better

resource utilization. Furthermore, the Skymont E-cores feature wider allocation and retire mechanisms, enabling E-core efficiency to match P-core performance levels. The Enhanced Memory subsystem, with 4x 128-bit FP & SIMD vector processing, claims to deliver double the AI throughput compared to the previous Meteor Lake platform. On the performance cluster, featuring four Lion Cove P-cores, Intel asserts the fastest single-core performance in the industry for thin and light AI PCs, within a power envelope of 8W-37W, with a maximum of 37W. Intel's Lion Cove P-core architecture offers enhanced processing efficiency through improved scheduling capabilities and 18 execution ports, alongside access to up to 12MB of shared L3 cache for accelerated AI and compute-intensive tasks. Additionally, Intel highlights its AI-driven power management system and enhanced memory subsystem as key optimizations for power per area (PPA), with finer control over power consumption via 16.67MHz clock interval control. According to Intel, these advancements achieve a balance between performance and efficiency in its P-cores.

Fastest integrated graphics for thin and light PCs?

Does Intel's Xe2 GPU Architecture Claim the Title of Fastest Integrated Graphics for Thin and Light PCs? Indeed, it does. The new Intel Xe2 GPU architecture boasts the fastest performance numbers for integrated GPUs in thin and light PCs. The Intel Arc 140V GPU features 8 cores with a maximum clock speed of 2.05 GHz, while the lower variant, the Arc 130V GPU, has 7 cores and reaches a maximum of 1.85 GHz. Intel claims an impressive 31% overall performance increase compared to the previous generation Xe1, which powered Meteor Lake-based laptops. During the performance briefing session, Intel showcased a gaming performance chart demonstrating thin and light laptops running on Intel Lunar